

## Homework 3 Written Answers

*Data Engineering 300*

*Aden Benson*

*5/26/2025*

### Part 1

#### Task 1

To begin Task 1, I first created a “words” Pyspark dataframe that exploded all document columns and listed each individual word with its corresponding document number. This is shown in figure 1.

+---+-----+	
_c0	word
+---+-----+	
0	wall
0	st
0	bears
0	claw
0	back
0	black
0	reuters
0	reuters
0	short
0	sellers
0	wall
0	street
0	dwindling
0	band
0	ultra
0	cynics
0	seeing
0	green
1	carlyle
1	looks
+---+-----+	

Fig. 1 - Exploded word dataframe

After this, I was able to design map/reduce functions to create dataframes for the following statements: For each  $d$ , the counts of  $t$ , for each  $d$  the counts of words, and for each  $t$ , the counts of  $d$  that contain  $t$ . The structure for the second and third dataframe are shown below in figures 2 and 3. The first dataframe was too large to use `.collect()`.

```
[(0, 18.0), (9, 19.0), (18, 21.0), (27, 21.0), (36, 26.0)]
```

Fig. 2 - For each  $d$  the counts of words (Doc number, count of words)

```
[('nation', 1598.0),
 ('eggs', 38.0),
 ('software', 3798.0),
 ('different', 410.0),
 ('7', 2614.0)]
```

Fig. 3 - For each  $t$ , the counts of  $d$  that contain  $t$  (Word, number of occurrences in documents)

## Task 2

In Task 2, we create functions that first calculate the TF term, then the IDF term, and combine these measures to compute the TF-IDF measure for each word in each document. In figure 4, we see an example of how this is displayed in our final dataframe.

```
+-----+-----+-----+
|doc_Number|  word|   tf_idf_measure|
+-----+-----+-----+
|      18|deficit|  0.540640233790213|
|    11970|deficit|0.24681401977379286|
|    20826|deficit|0.23653010228321816|
|    21078|deficit|0.23170295733866267|
|    37116|deficit|0.24681401977379286|
+-----+-----+-----+
```

Fig. 4 - Final dataframe that contains a document number, a word, and its TF-IDF measure.

## Task 3

In Task 3, we find the TF-IDF measure for each word in the first five documents. These had to be displayed separately since there are over 10 words in each document. In figures 5-6, we see these measures.

doc_Number	word	tf_idf_measure	doc_Number	word	tf_idf_measure	doc_Number	word	tf_idf_measure
	0	seeing 0.37743394553516213		1	placed 0.2284965552404658		2	expected 0.16094627131903613
	0	short 0.2773120373951269		1	industry 0.15043731768548949		2	earnings 0.1792714404894228
	0	claw 0.499114829314058		1	aerospace 0.2581171817448437		2	stock 0.17879168082328206
	0	back 0.1892216338539946		1	quietly 0.25188254045524316		2	prices 0.14472559202114177
	0	green 0.2877107940095433		1	looks 0.1973537176743789		2	hang 0.30475018305843793
	0	dwindling 0.4572386180709258		1	bets 0.27861293130724324		2	worries 0.23009353850726894
	0	ultra 0.4125512394225831		1	toward 0.1898997183872362		2	stocks 0.14976769101715193
	0	st 0.2584728642725166		1	carlyle 0.2578098186776328		2	depth 0.31343954772064864
	0	sellers 0.4468379768438066		1	firm 0.15969712503706046		2	outlook 0.4265073217271922
	0	street 0.24678348986493034		1	defense 0.1751279339938823		2	oil 0.13908157105107033
	0	cynics 0.563734318747707		1	plays 0.22418048797172685		2	summer 0.22694739048609625
	0	bears 0.3372044607529448		1	investment 0.1890771769001148		2	market 0.15069298739291276
	0	wall 0.5115985326511431		1	commercial 0.2057832028092643		2	next 0.14062721303262238
	0	reuters 0.24754017186645658		1	market 0.13394932212703356		2	soaring 0.2596334462817101
	0	black 0.2953171727366614		1	private 0.1929050573011279		2	crude 0.197241148492091
	0	band 0.3643421454792778		1	reputation 0.2578098186776328		2	economy 0.3721400726458204
				1	timed 0.324478643568105		2	week 0.13121900794126834
				1	making 0.1698717076460444		2	plus 0.24449073714833106
				1	group 0.12468100563149095		2	cloud 0.295159450642955
				1	part 0.16022031730914288		2	reuters 0.18565512889984243

Fig. 5 - TF-IDF measures for documents 1-3

doc_Number	word	tf_idf_measure	doc_Number	word	tf_idf_measure
	3	iraq 0.23809526243476142		4	time 0.10623532598945136
	3	showed 0.1743365558077232		4	us 0.1669859687392097
	3	intelligence 0.20782569445751425		4	present 0.22209684830286883
	3	strike 0.17411586950893898		4	record 0.1232987151692413
	3	official 0.15149485319300557		4	months 0.14002501854271598
	3	flows 0.2774168429760197		4	soar 0.2306791247647116
	3	main 0.36492623402353547		4	prices 0.23156094723382684
	3	said 0.06593367258642661		4	barely 0.21935019724396657
	3	rebel 0.18209445014364567		4	wallets 0.2665151844733088
	3	militia 0.2252006141545402		4	menace 0.5747440955975784
	3	oil 0.35763832555989516		4	toppling 0.27964532733021175
	3	pipeline 0.4720829409342409		4	oil 0.22253051368171256
	3	infrastructure 0.22959926718225876		4	tearaway 0.3918885216630942
	3	export 0.23862435123782139		4	new 0.1271397626254836
	3	authorities 0.18159366801541998		4	straining 0.2904044404056468
	3	southern 0.336553609483104		4	economy 0.14885602905832815
	3	saturday 0.12197305137253434		4	posing 0.2589223867776184
	3	halts 0.27365396741681164		4	economic 0.14782686453681568
	3	halted 0.2557691357056513		4	world 0.09332201126546583
	3	reuters 0.15913296762843637		4	elections 0.16009904796740967

Fig. 6 - TF-IDF measures for documents 4-5

## Part 2

### Task 1

In Task 1, I first designed MapReduce functions that would later be used in the loss\_SVM() function to calculate a sum for the loss function. The summed total is shown below in figure 7.

```
# Design reduce portion, which finds the sum of all  $\max(0, 1 - y_i(w^T x_i + b))$  values
from operator import add
summed_total = svm_ps.rdd.map(map_part2).reduceByKey(add).collect()[0][1]
summed_total
```

✓ 2.4s

399889.5049647091

Fig. 7 - Summed total used below in loss\_SVM function

## Task 2

Task 2 has no visible outputs, as it is simply defining a function to compute the loss objective function given weights, bias, X, y, and a lambda value.

## Task 3

In Task 3, we compute the loss\_SVM value given our data uploaded into Pyspark. We see that the objective value for the given data is 1.0029404.

```
# Create X and y dfs
X = svm_ps.select(svm_ps.columns[:-1])
y = svm_ps.select(svm_ps.columns[-1])

print('Objective value: ', loss_SVM(w_ps, bias_ps, X, y, 1))
```

✓ 7.5s

[Stage 85:=====>

(1 + 9) / 10]

Objective value: 1.0029403834857487

Fig. 8 - Loss objective value from given data

## Task 4

Finally, we design MapReduce functions to make predictions using the  $\hat{y}$  function provided in the assignment documentation. In figure 9, we see the structure of these predictions. In figure 10, we find that the loss objective value for our  $\hat{y}$  predictions is 0.9579.

```
[('Prediction', -1),
 ('Prediction', -1),
 ('Prediction', -1),
 ('Prediction', 1),
 ('Prediction', -1),
 ('Prediction', 1),
 ('Prediction', -1),
 ('Prediction', -1),
 ('Prediction', 1),
 ('Prediction', -1),
 ('Prediction', 1),
 ('Prediction', -1),
 ('Prediction', -1),
 ('Prediction', 1),
 ('Prediction', 1),
 ('Prediction', 1),
 ('Prediction', -1),
 ('Prediction', 1),
 ('Prediction', -1),
 ('Prediction', 1),
 ('Prediction', 1),
 ('Prediction', 1),
 ('Prediction', -1),
 ('Prediction', -1),
 ...]
```

Fig. 9 - Structure of prediction dataframe.

```
# For comparison, find the loss of our predictions
preds_df=preds.toDF().select(col('_2').alias('_c64'))
print('Loss of Predictions: ', loss_SVM(w_ps, bias_ps, X, preds_df,1))
```

✓ 13.2s

[Stage 101:====>

(1 + 9) / 10]

Loss of Predictions: 0.9579132416284583

Fig. 10 - Loss objective value for  $\hat{y}$  predictions.